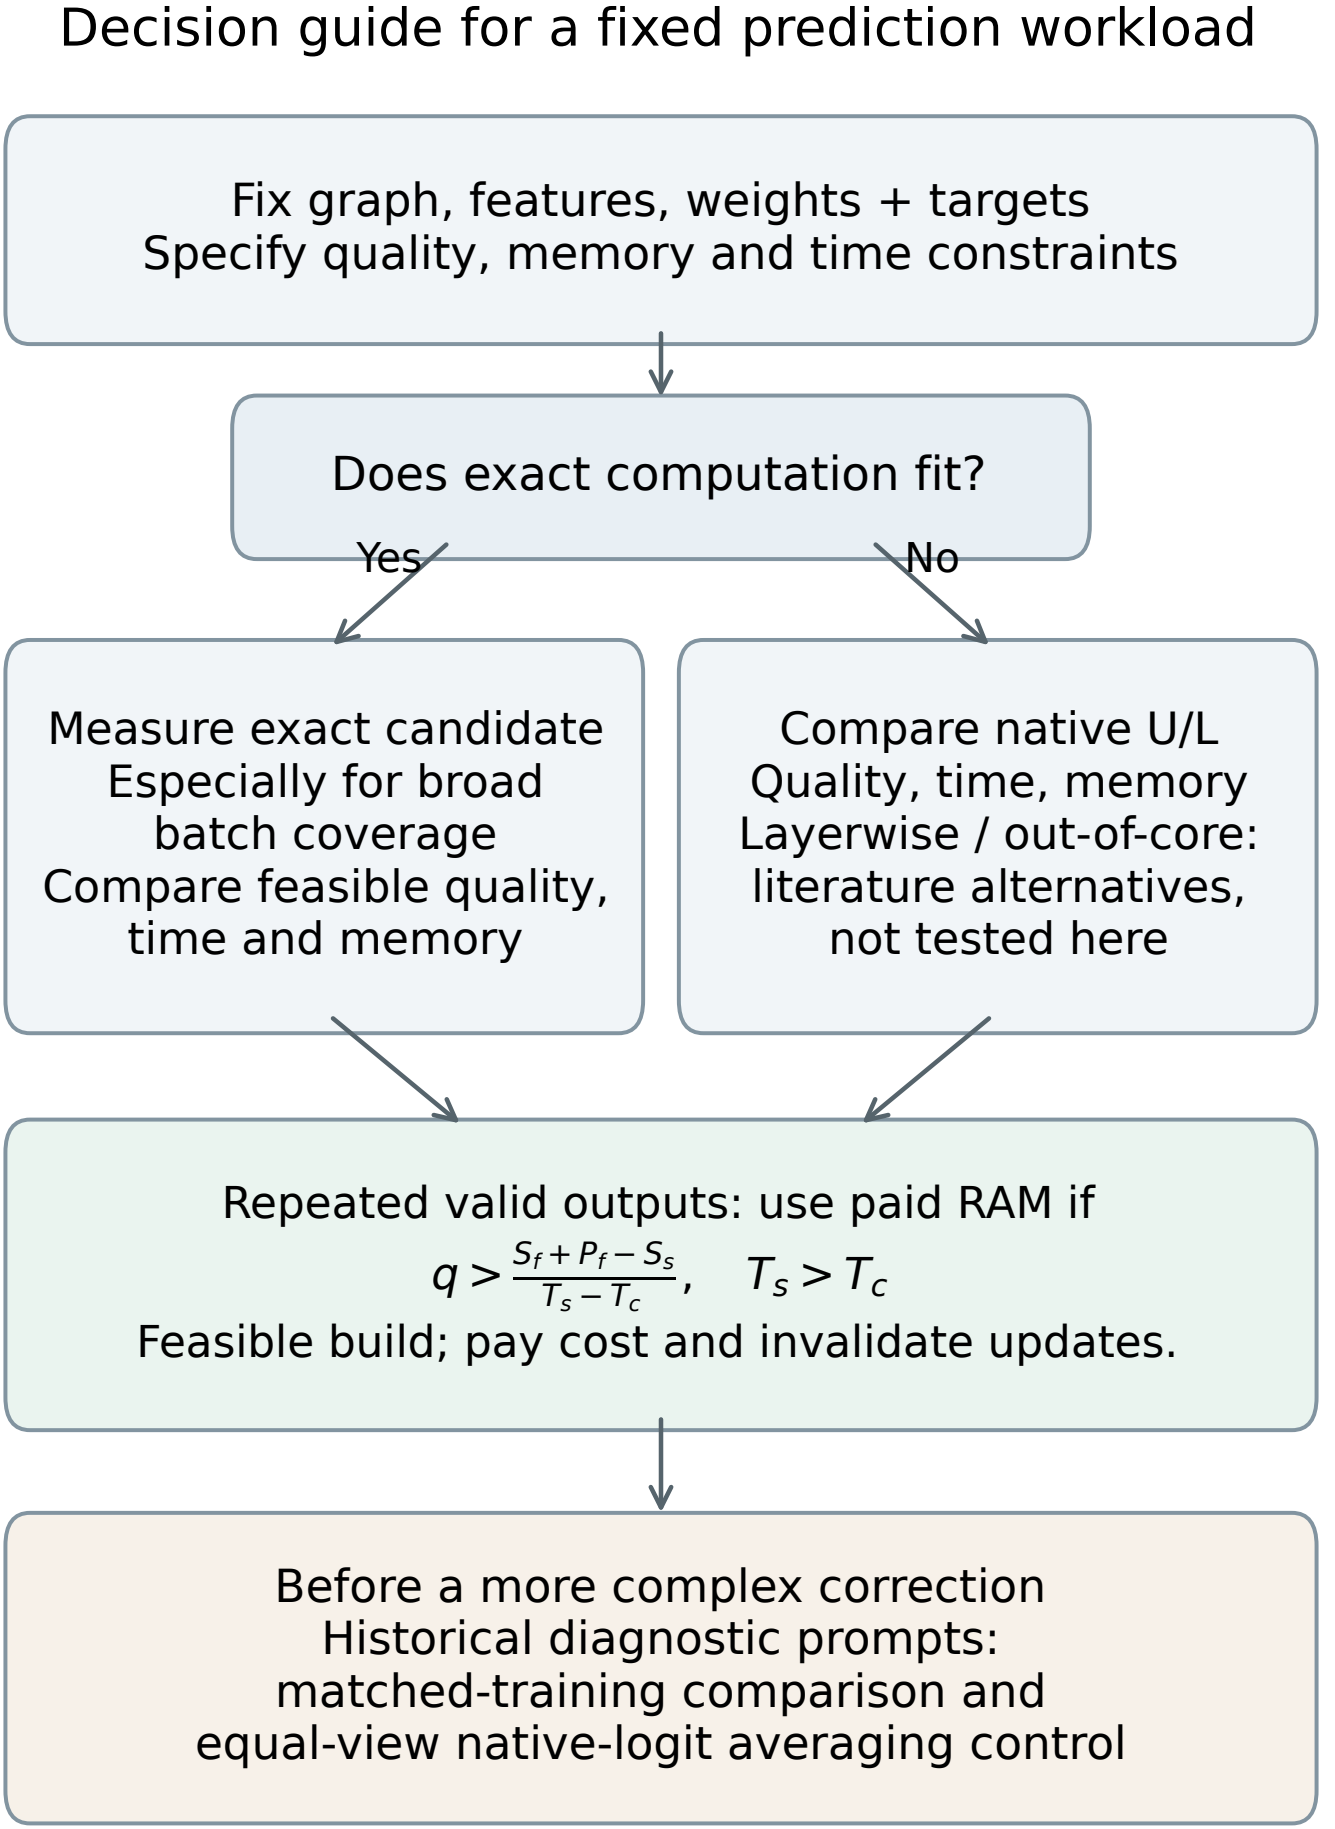
Decision guide for a fixed prediction workload
Fix graph, features, weights + targets
Specify quality, memory and time constraints
Does exact computation fit?
Yes
No
Measure exact candidate
Especially for broad
batch coverage
Compare feasible quality,
time and memory
Compare native U/L
Quality, time, memory
Layerwise / out-of-core:
literature alternatives,
not tested here
Repeated valid outputs: use paid RAM if
$q > \frac{S_f + P_f - S_s}{T_s - T_c}, \quad T_s > T_c$
Feasible build; pay cost and invalidate updates.
Before a more complex correction
Historical diagnostic prompts:
matched-training comparison and
equal-view native-logit averaging control
Compare complete targets; keep cache versions valid.