







Existing bounded runs: learning and layer-specific radial concentration

Seed 11



Seed 23



Descriptive existing logs only. Near cap means scaled message radius $\geq 0.95 \times 1.2$; not manifold numerical failure. Training loss and validation retrieval are different objectives.